

Data Challenge: Exploratory Data Analysis

Based on material developed by Sam Clifford

Introduction

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



Suite of R packages to make working with data as easy as possible (Wickham 2020), including

- ggplot2: for plotting data
- dplyr: for manipulating data frames
- tidyr: for making data tidy
- forcats: for manipulating factor variables
- magrittr: for easy chaining of commands

(Wickham and Grolemund 2017; Wickham et al. 2019)

Summarising Data

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



- Summary statistics one of most common data analysis tasks
- Consider the `Gestation` data from `mosaicData`
- Birth weight, date, and gestational period collected as part of the Child Health and Development Studies in 1961 and 1962. Information about the baby's parents — age, education, height, weight, and whether the mother smoked is also *recorded* (Nolan and Speed 2001).
- We will use some functions from `dplyr` to choose, group and summarise data
- `verb(.data, ...)` applies a `dplyr verb` to a data frame

Summarising Data

- `count` how many babies in data set

```
library(mosaicData)
library(tidyverse)
data(Gestation)

count(Gestation) # just like nrow()
## # A tibble: 1 × 1
##       n
##   <int>
## 1  1236
```

Summarising Data

- Can also count for a given grouping variable

```
count(Gestation, race) # nrow can't do this
## # A tibble: 6 × 2
##   race      n
##   <chr> <int>
## 1 asian    44
## 2 black   244
## 3 mex     40
## 4 mixed    25
## 5 white   870
## 6 <NA>    13
```

Summarising Data

- The `summarise` function allows us to calculate summary statistics of a variable
- Can (and should) give names to summary columns
- Calculate the mean birth weight in the data set

```
summarise(Gestation, wt_mean = mean(wt))  
## # A tibble: 1 × 1  
##   wt_mean  
##   <dbl>  
## 1    120.
```


Summarising Data

- We can calculate multiple summaries at once

```
summarise(Gestation,  
          Mean = mean(wt),  
          SD   = sd(wt),  
          Low  = quantile(wt, 0.025),  
          High = quantile(wt, 0.975))  
## # A tibble: 1 × 4  
##   Mean    SD   Low  High  
##   <dbl> <dbl> <dbl> <dbl>  
## 1  120.  18.2    81   155
```

- summarise() applies **summary functions** to columns to create a new table.
- Summary functions **take** vectors as input and **return single values** as output.

■ ■ ■ **summary function** → ■

Summarising grouped data

- We can group the rows in our data and calculate summaries for each group
- `group_by` lets us pass variable names to set the structure
- Row order is maintained

```
Gestation_grouped_by_race <- group_by(Gestation, race)
```

```
Gestation_grouped_by_race
```

```
## # A tibble: 1,236 × 23
```

```
## # Groups:   race [6]
```

```
##      id plurality outcome date      gestation sex      wt parity race      age
##      <dbl> <chr>      <chr> <date>      <dbl> <chr> <dbl> <dbl> <chr> <dbl>
##  1     15 single fe... live bi... 1964-11-11      284 male    120      1 asian    27
##  2     20 single fe... live bi... 1965-02-07      282 male    113      2 white    33
##  3     58 single fe... live bi... 1965-04-25      279 male    128      1 white    28
##  4     61 single fe... live bi... 1965-02-12      NA male    123      2 white    36
##  5     72 single fe... live bi... 1964-11-25      282 male    108      1 white    23
##  6    100 single fe... live bi... 1965-07-31      286 male    136      4 white    25
##  7    102 single fe... live bi... 1964-12-19      244 male    138      4 black    33
##  8    129 single fe... live bi... 1965-04-11      245 male    132      2 black    23
##  9    142 single fe... live bi... 1964-11-08      289 male    120      3 white    25
## 10    148 single fe... live bi... 1965-04-17      299 male    143      3 white    30
## # ... with 1,226 more rows, and 13 more variables: ed <chr>, ht <dbl>,
## # wt.1 <dbl>, drace <chr>, dage <dbl>, ded <chr>, dht <dbl>, dwt <dbl>,
## # marital <chr>, inc <chr>, smoke <chr>, time <chr>, number <chr>
```

Summarising grouped data

- `summarise()` respects the grouping structure

```
summarise(Gestation_grouped_by_race,  
          Mean = mean(wt),  
          SD   = sd(wt),  
          Low  = quantile(wt, 0.025),  
          High = quantile(wt, 0.975))  
## # A tibble: 6 × 5  
##   race    Mean    SD    Low    High  
##   <chr> <dbl> <dbl> <dbl> <dbl>  
## 1 asian  110.   16.0  78.4  139.  
## 2 black  113.   19.1  71    150  
## 3 mex    124.   14.1  99.0  146.  
## 4 mixed  120.   20.1  78.8  150.  
## 5 white  122.   17.7  85    158  
## 6 <NA>  117.   16.7  86.8  143.
```

Rows and columns

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



Operating on columns

- create/modify/delete columns with dplyr's `mutate()`
- e.g. relabelling `race` so words start with a capital,

```
Gestation <- mutate(Gestation, race = str_to_title(race))
```

```
count(Gestation, race)
```

```
## # A tibble: 6 × 2
```

```
##   race      n
```

```
##   <chr> <int>
```

```
## 1 Asian    44
```

```
## 2 Black   244
```


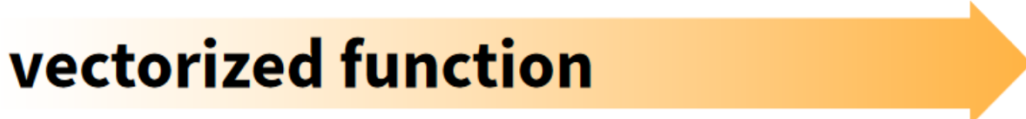

```
## 3 Mex     40
```

```
## 4 Mixed   25
```

```
## 5 White  870
```

```
## 6 <NA>    13
```

- `mutate()` applies **vectorized functions** to columns to create new columns.
- Vectorized functions **take** vectors as input and return vectors of the **same length** as output.

 **vectorized function**  

Choosing columns

- For one reason or another we may want to select only certain columns of our data frame

```
head(Gestation, 1)
## # A tibble: 1 × 23
##       id plurality    outcome date      gestation sex      wt parity race    age
##   <dbl> <chr>      <chr>   <date>      <dbl> <chr> <dbl>  <dbl> <chr> <dbl>
## 1    15 single fetus live b... 1964-11-11    284 male   120      1 Asian    27
## # ... with 13 more variables: ed <chr>, ht <dbl>, wt.1 <dbl>, drace <chr>,
## #   dage <dbl>, ded <chr>, dht <dbl>, dwt <dbl>, marital <chr>, inc <chr>,
## #   smoke <chr>, time <chr>, number <chr>
```

```
head(select(Gestation, race, wt, number), 1)
## # A tibble: 1 × 3
##   race      wt number
##   <chr> <dbl> <chr>
## 1 Asian   120 never
```


Choosing and renaming columns

- We can also rename columns on the fly as we select them

```
select(Gestation,  
       Race           = race,  
       `Birthweight (oz)` = wt,  
       `Cigs. smoked`   = number)  
## # A tibble: 1,236 × 3  
##   Race `Birthweight (oz)` `Cigs. smoked`  
##   <chr>          <dbl> <chr>  
## 1 Asian             120 never  
## 2 White             113 never  
## 3 White             128 1-4 per day  
## 4 White             123 20-29 per day  
## 5 White             108 20-29 per day  
## 6 White             136 5-9 per day  
## 7 Black             138 never  
## 8 Black             132 never  
## 9 White             120 never  
## 10 White            143 15-19 per day  
## # ... with 1,226 more rows
```

Choosing and renaming columns

- Alternatively we can rename columns without worrying about failing to select columns we haven't renamed

```
names(Gestation)
## [1] "id"          "plurality"   "outcome"    "date"       "gestation"  "sex"
## [7] "wt"          "parity"     "race"       "age"        "ed"         "ht"
## [13] "wt.1"       "drace"      "dage"       "ded"        "dht"        "dwt"
## [19] "marital"    "inc"        "smoke"      "time"       "number"
```

```
Gestation <- rename(Gestation, `Cigs. smoked` = number)
```

```
names(Gestation)
## [1] "id"          "plurality"   "outcome"    "date"       "gestation"
## [6] "sex"         "wt"          "parity"     "race"       "age"
## [11] "ed"          "ht"          "wt.1"       "drace"      "dage"
## [16] "ded"         "dht"         "dwt"        "marital"    "inc"
## [21] "smoke"       "time"        "Cigs. smoked"
```


Choosing rows

- The dplyr equivalent of subset is filter
- Takes a logical statement and does non-standard evaluation of variable names •
`filter(data, A & B)` the same as `filter(data, A, B)`

```
Gestation2 <- select(Gestation,
                     Race           = race,
                     `Birthweight (oz)` = wt,
                     `Cigs. smoked`)

filter(Gestation2, Race == "White", `Cigs. smoked` == "never")
## # A tibble: 352 × 3
##   Race `Birthweight (oz)` `Cigs. smoked`
##   <chr>          <dbl> <chr>
## 1 White             113 never
## 2 White             120 never
## 3 White             144 never
## 4 White             125 never
## 5 White             122 never
## 6 White             113 never
## 7 White             134 never
## 8 White             128 never
## 9 White             129 never
## 10 White            110 never
## # ... with 342 more rows
```

Choosing rows

- `slice*()` functions allow you to select rows based on their properties, e.g. which babies have lowest birth weight overall and in each race group?

```
slice_min(Gestation2, `Birthweight (oz)`)
## # A tibble: 1 × 3
##   Race `Birthweight (oz)` `Cigs. smoked`
##   <chr>          <dbl> <chr>
## 1 Black              55 never

slice_min(group_by(Gestation2, Race), `Birthweight (oz)`)
## # A tibble: 6 × 3
## # Groups:   Race [6]
##   Race `Birthweight (oz)` `Cigs. smoked`
##   <chr>          <dbl> <chr>
## 1 Asian              71 5-9 per day
## 2 Black              55 never
## 3 Mex                97 never
## 4 Mixed              77 20-29 per day
## 5 White              63 never
## 6 <NA>              82 20-29 per day
```

Reshaping data frames

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



“**TIDY DATA** is a standard way of mapping the meaning of a dataset to its structure.”

—HADLEY WICKHAM

In tidy data:

- each variable forms a column
- each observation forms a row
- each cell is a single measurement

each column a variable

id	name	color
1	floof	gray
2	max	black
3	cat	orange
4	donut	gray
5	merlin	black
6	panda	calico

each row an observation

Long and wide tidy data

To make this pivot, we specify

- which `cols` are to be converted from being columns of length k to one column of length $n \times k$
- the `names_to` column name, which contains the names of the pivoted columns
- the `values_to` name of the column containing the `value` of each variable for each `id`

wide

id	x	y	z
1	a	c	e
2	b	d	f

Long and wide tidy data

```
Gestation_igwa <- select(Gestation, id, gestation, wt, age)
```

```
Gestation_long <- pivot_longer(  
  data      = Gestation_igwa,  
  cols      = c(gestation, wt, age),  
  names_to  = 'name',  
  values_to = 'value')
```

```
head(Gestation_long, 6)  
## # A tibble: 6 × 3  
##       id name      value  
##   <dbl> <chr>    <dbl>  
## 1    15 gestation    284  
## 2    15 wt          120  
## 3    15 age           27  
## 4    20 gestation    282  
## 5    20 wt          113  
## 6    20 age           33
```

NB: we need to use ' quotes for names_to and values_to arguments because they are strings defining new columns

Long and wide tidy data

Or specify which columns *not* to pivot, e.g. `-id` selects all variables except `id`

```
Gestation_long <- pivot_longer(Gestation_igwa, -id)
```

```
head(Gestation_long, 6)
```

```
## # A tibble: 6 × 3
```

```
##       id name      value
```

```
##   <dbl> <chr>    <dbl>
```

```
## 1     15 gestation  284
```

```
## 2     15 wt         120
```

```
## 3     15 age         27
```

```
## 4     20 gestation  282
```

```
## 5     20 wt         113
```

```
## 6     20 age         33
```


Long and wide tidy data

- To convert to a wider format, we use `pivot_wider`
- For example, we specify:
- The *data* source
 - where we get the new column *names from*
 - where we get the new column *values from*

```
Gestation_wide <- pivot_wider(data      = Gestation_long,  
                             names_from = name,  
                             values_from = value)
```

```
head(Gestation_wide, 3) # recovered original data frame
```

```
## # A tibble: 3 × 4
```

```
##       id gestation    wt  age
```

```
##   <dbl>    <dbl> <dbl> <dbl>
```

```
## 1     15      284   120   27
```

```
## 2     20      282   113   33
```

```
## 3     58      279   128   28
```


Pipe

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



- dplyr imports the `%>%` pipe from magrittr
- `f (g (x))` is equivalent to `x %>% g %>% f`
- Makes it easier to chain operations together without storing temporary objects
- Output on left of `%>%` becomes first argument of function on right
 - by convention, all tidyverse functions take a data frame as their first argument

```
x %>% f_1 %>% f_2 %>% f_3
```

```
# rather than
```

```
f_3(f_2(f_1(x)))
```

```
# or even worse...
```

```
x_1 <- f_1(x)
```

```
x_2 <- f_2(x_1)
```

```
x_3 <- f_3(x_2)
```

- An example

```
Gestation %>% group_by(race, smoke) %>%  
  summarise(wt = mean(wt)) %>% pivot_wider(names_from = race,  
                                             values_from = wt)
```

A tibble: 5 × 7

##	smoke	Asian	Black	Mex	Mixed	White	`NA`
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
##	1 never	114.	117.	123.	125.	125.	119.
##	2 now	98.8	108.	127.	108.	116.	114.
##	3 once did, not now	111	117.	NA	134	126.	NA
##	4 until current pregnancy	117	113.	129.	111	127.	NA
##	5 <NA>	126	119	115	NA	131.	NA

- NB the `smoke` variable is character and sorted alphabetically.
- We don't *expect* you to use pipes, but they're useful

Summary

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



- Summarising data
 - `group_by` to set a group structure
 - `summarise` to calculate summary stats across group structure
 - `count` to see how many rows in each group
- Reshaping data frames
 - `pivot_longer` from variables side by side to key-value
 - `pivot_wider` from key-value to named column variables

- Dealing with rows and columns
 - `mutate` to create/modify/delete columns
 - `select` to choose columns
 - `filter` to choose rows based on logical condition
 - `slice*` to choose rows based on position or property
- Pipe
 - `%>%` to chain operations
- Wickham (2014) on what tidy data is
- Wickham et al. (2019) for more explanation of tidyverse

Visualisation with ggplot

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



Why do we visualise?

Since the aim of exploratory data analysis is to learn what seems to be, it should be no surprise that pictures play a vital role in doing it well. There is nothing better than a picture for making you think of questions you had forgotten to ask (even mentally).

Tukey and Tukey (1985)

Tufte (1983) and Pantoliano (2012)

- Show the data
- Provide clarity
- Allow comparison where appropriate
 - use aesthetics to draw attention to important details
 - make clear that data has multiple levels of structure

- Produce graphs with high data density
 - make every drop of ink count
 - careful use of whitespace
- Avoid excessive and unnecessary use of graphical effects
- Reader should be able to understand what the graph means and not be
 - misled into thinking something that is untrue
 - distracted from the main point

`ggplot2` uses a grammar of graphics (Wickham 2010)

- map variables in data frame to aesthetic options in the plot
- choose a geometry for how to display these variables
- adjustments to axis scales
- adjustments to colors, themes, etc.
- adding extra commands in a ‘do this, then do this’ manner
- python users have `plotnine` (Kibirige 2020) which is based on the same ideas

How do we structure a call to `ggplot` to make a plot?

```
# library(ggplot2) # already loaded with tidyverse  
ggplot(data = my.data.frame,  
       aes(x = my.x.variable,  
           y = my.y.variable)) +  
  geom_point()
```


- Load `ggplot2` package
- Specify we want a `ggplot` object and which data frame we're going to use,
- Set **aesthetic options** to map to the axes of the plot
- State geometry we're using to show variables

Building plots

- For example, consider daily maximum temperature varying with solar radiation in New York City 1973
- Each row is a pair of values (x, y), shown as a point

```
data(airquality)
solar_temp_plot <- ggplot(data = airquality,
  aes(x = Solar.R, y = Temp)) + geom_point()


solar_temp_plot
```



Scatter plot

We can add some human-friendly labels and change the theme

```
solar_temp_plot <- solar_temp_plot + theme_bw() +  
  labs(x = 'Solar radiation (Langleys)', y = 'Max. daily temp. (F)')  
  
solar_temp_plot
```



Some more geometries

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



Line plot


- Similar to scatter plot, but joins pairs of values
- Useful when showing how something changes over time
- If (x, y) pairs ordered by
 - x, use `geom_line()` (e.g. x is time)
 - row order, use `geom_path()`
 - nothing, don't use a line

```
airquality <- mutate(airquality, Date = as.Date(paste('1973', Month, Day, sep = '-')))  
  
airquality_plot <- ggplot(data = airquality, aes(x=Date, y=Ozone)) +  
  geom_line() + theme_bw() + labs(y = 'Ozone concentration (ppb)',  
                                   title = 'Daily mean Ozone in NYC (1973)')
```


Line plot

```
airquality_plot
```

Daily mean Ozone in NYC (1973)




- You may see this referred to as a time series plot

Line plot

- Observations whose neighbours are NA values can't be plotted with a line
- Can layer multiple geometries for same aesthetic mapping

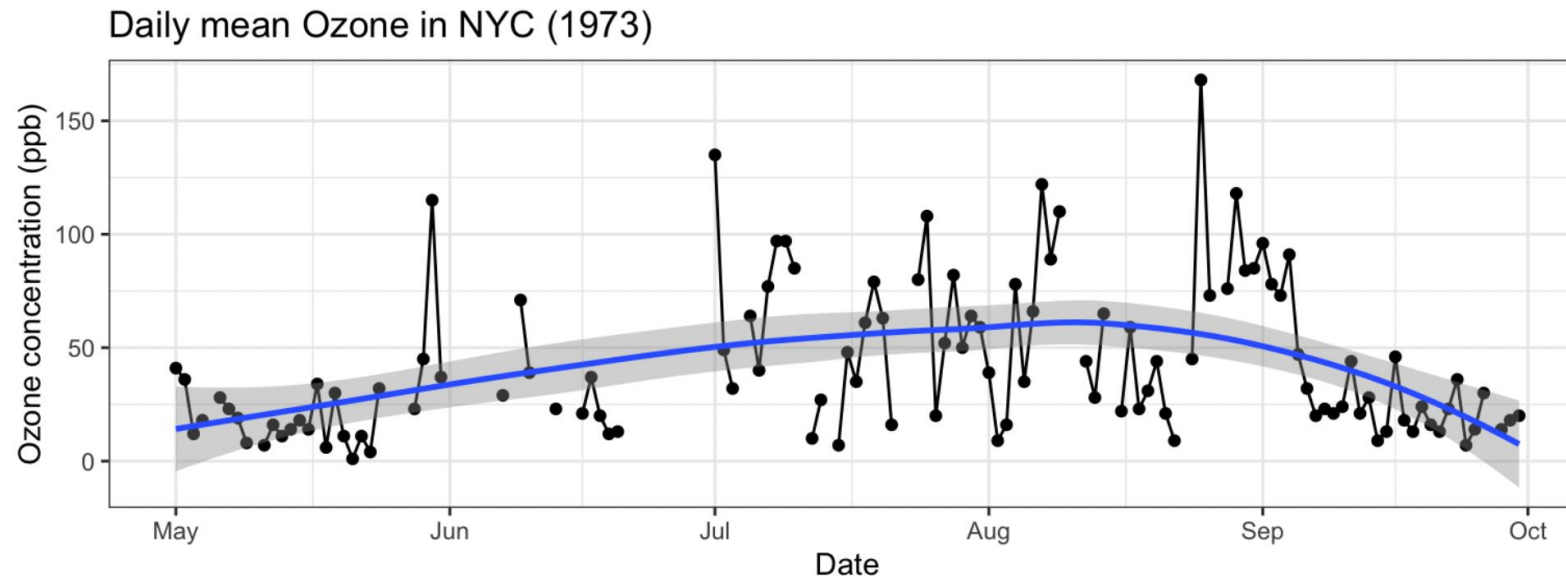
```
airquality_plot + geom_point()
```



Scatterplot smoother

- Often too much data in a scatter plot to see pattern
- Maybe we want to highlight the trend in the data


```
airquality_plot + geom_point() + geom_smooth()
```



Boxplot

- continuous y, discrete x
- outliers (> 1.5 IQR from median) shown as points automatically

```
ggplot(data = airquality, aes(x = factor(Month), y = Ozone)) +  
  geom_boxplot() + theme_bw() +  
  labs(y = 'Ozone conc. (ppb)', x = 'Month')
```




Histograms

- univariate graphical summary needs only one aesthetic, `x`

```
ozone_hist <- ggplot(data = airquality, aes(x = Ozone)) +  
  geom_histogram(binwidth = 10, boundary = 0) +  
  labs(x = 'Ozone concentration (ppb)') + theme_bw()
```

ozone_hist




Density plots

- Kernel smoothing (continuous analogue of histogram)

```
ozone_dens <- ggplot(data = airquality, aes(x = Ozone)) +  
  geom_density(fill = 'grey35') +  
  labs(x = 'Ozone concentration (ppb)') + theme_bw()
```

ozone_dens



More on aesthetics

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



More on aesthetics

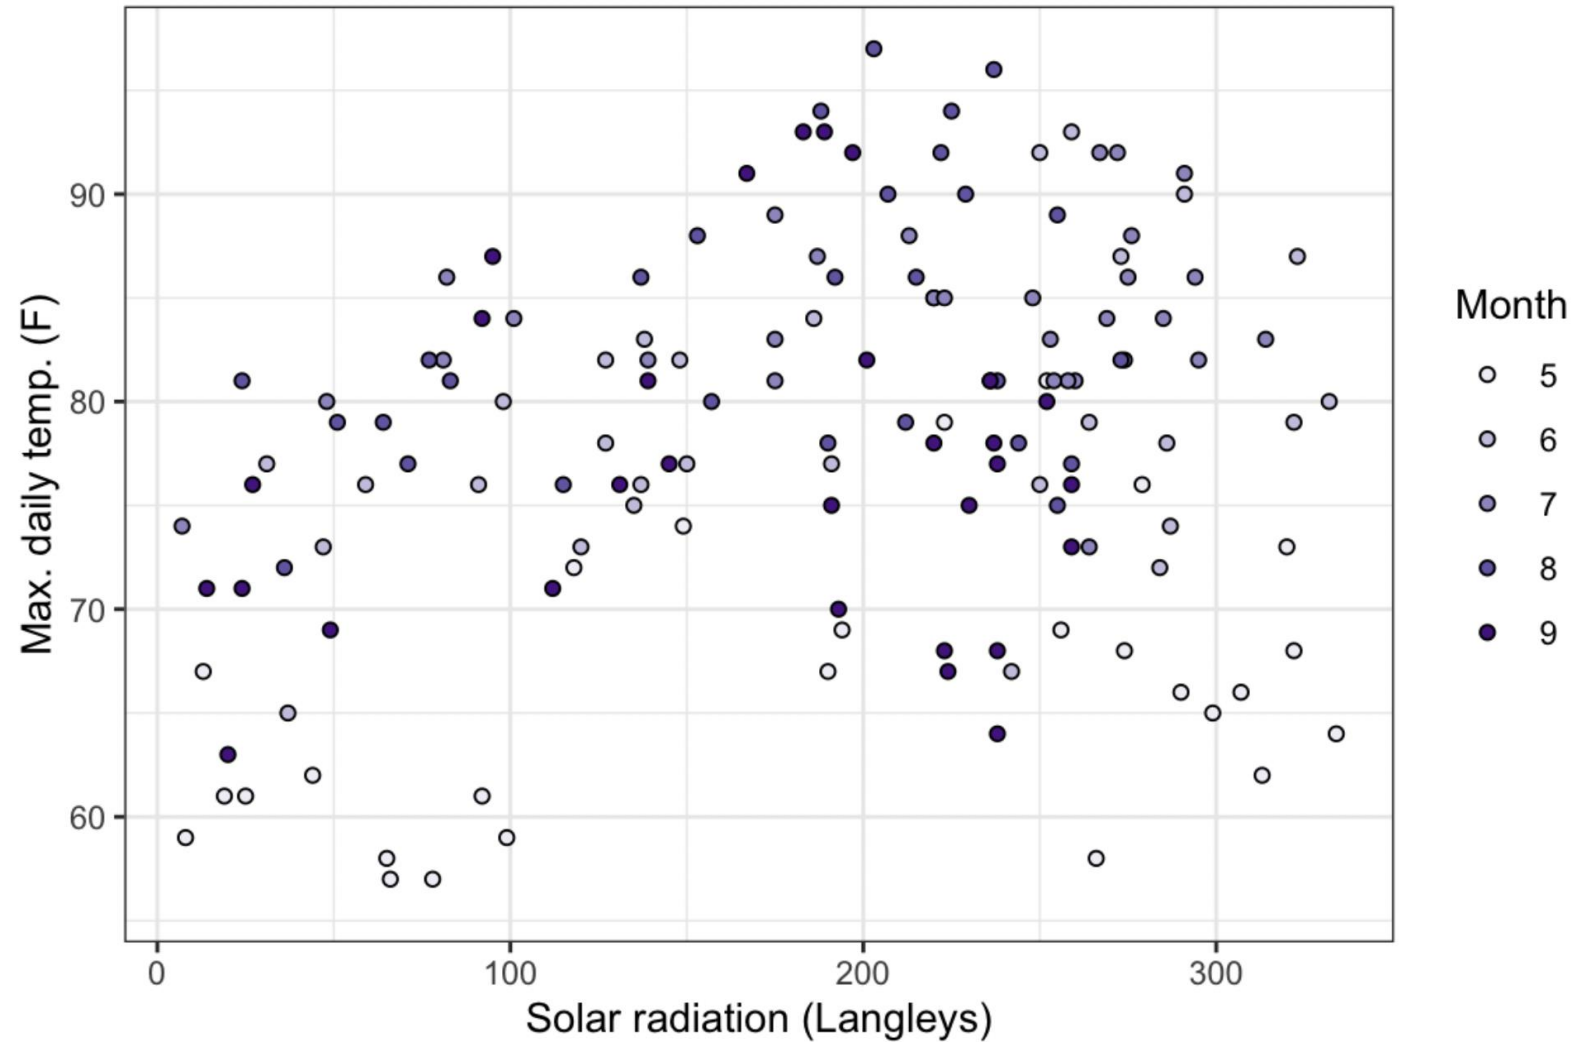
Aesthetic	What it effects
<code>size</code>	points, lines
<code>shape</code>	points
<code>linetype</code>	lines
<code>colour</code>	points, lines, boundary
<code>alpha</code>	transparency
<code>fill</code>	interior
<code>group</code>	repeats geometry

- If these (except `group`) are *outside* `aes()` they fix the value for all parts of that geometry
- Aesthetics specified inside `ggplot()` are inherited by all geometries for that plot
- Not all geometries accept all aesthetics (e.g. `geom_line()` has no fill)
- Some point shapes admit a `colour` and a `fill`

More on aesthetics

```
data(airquality)
solar_temp_plot_colored <-
  ggplot(data = airquality,
    aes(x = Solar.R, y = Temp)) +
  geom_point(aes(fill = factor(Month)),
    shape = 21,
    color = 'black') +
  labs(x = 'Solar radiation (Langleys)',
    y = 'Max. daily temp. (F)') +
  theme_bw() +
  scale_fill_brewer(palette = "Purples",
    name = 'Month')
```

More on aesthetics



Small multiples

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



Small multiples

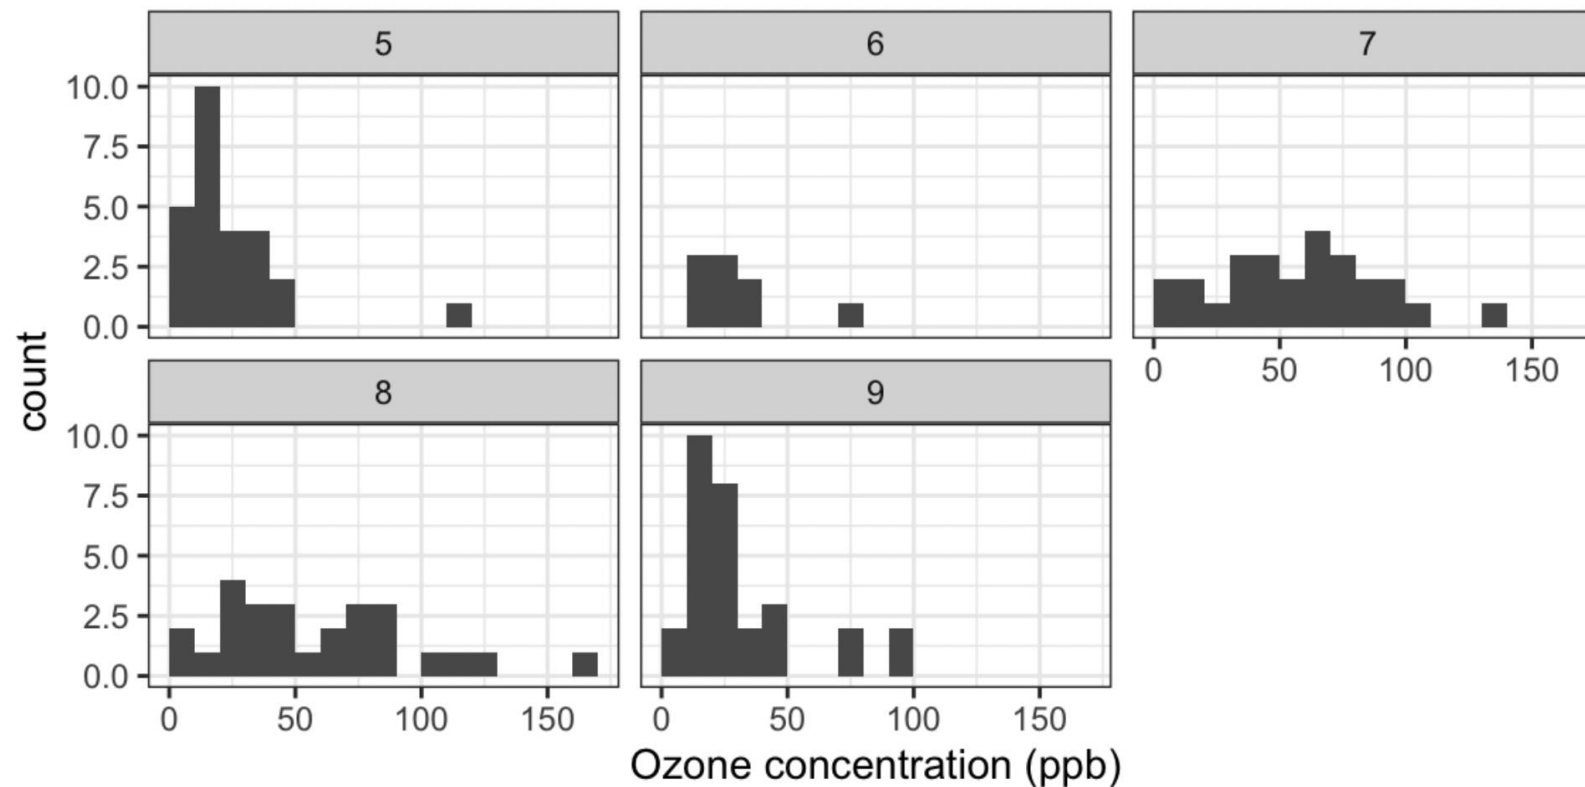
- Group a plot by some categorical variable
- Repeat a basic graph for groups in the data
 - air quality data has information about, e.g. months
- Can view 3-5 dimensions in the data on a 2D page
 - Often a better alternative to 3D, since it doesn't distort comparisons
 - Inner axes relate to the smallest X-Y plots
 - Outer axes relate to the grouping variables
- Avoids writing loops

Small multiples



- Repeat histogram for each value of Month, one per facet

```
ozone_hist + facet_wrap( ~ Month)
```



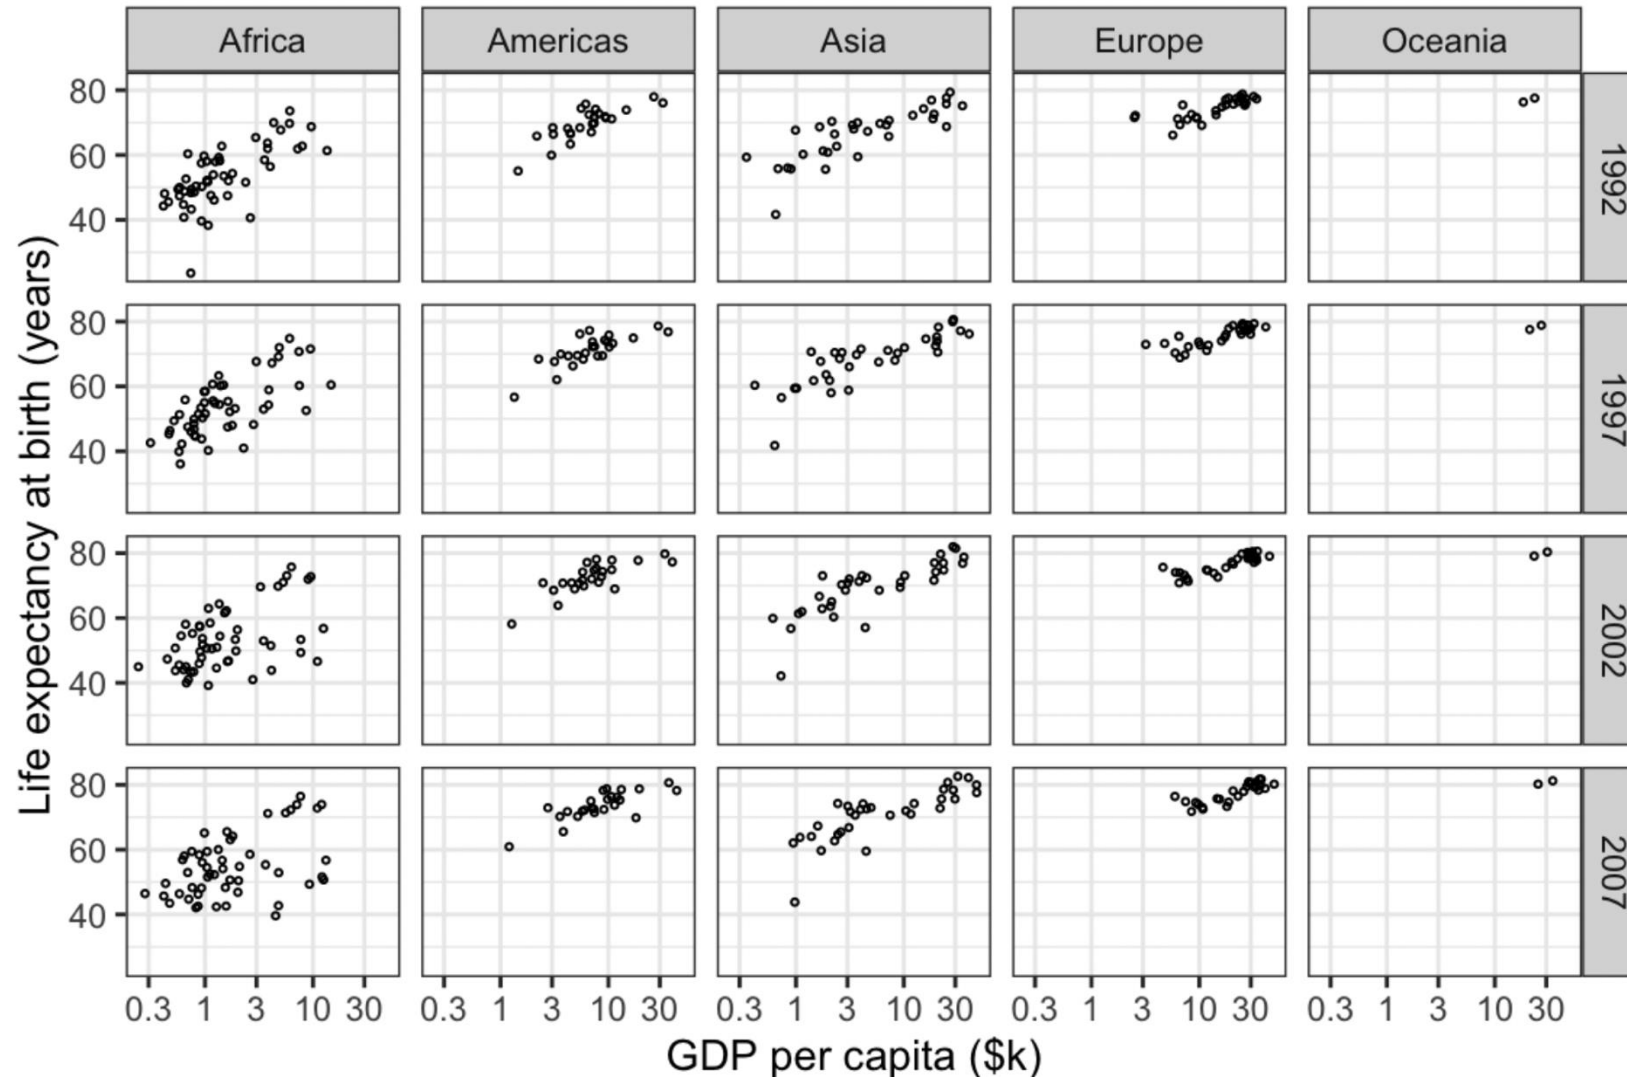
Small multiples

We can also use `facet_grid()` to repeat the aesthetic and geometries for specified `rows` and `cols` variables

```
library(gapminder)
data(gapminder)

gapminder_plot <-
  ggplot(data = subset(gapminder, year >= 1992),
    aes(x = gdpPercap/1e3,
      y = lifeExp)) +
  geom_point(shape = 1, size = 0.5) +
  facet_grid(rows = vars(year),
    cols = vars(continent)) +
  scale_x_log10(labels = ~sprintf("%g", .)) +
  xlab("GDP per capita ($k)") +
  ylab("Life expectancy at birth (years)") +
  theme_bw() +
  theme(panel.grid.minor.x = element_blank())
```

Small multiples



Summary

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



- We make graphs to tell a story with data
- Should draw reader in and explain what they're seeing
- Plots are built from
 - geometric objects
 - axis scales
 - coordinate systems (linear or logarithmic scale, 2D, 3D, etc.)
 - annotations (e.g. heading in small multiples)

- Successively building a plot with a grammar of graphics allows development of complex plots from simple elements and small changes
- Choose a plotting geometry that helps tell the story
- Meaningful labels remove ambiguity and confusion
- Be careful not to put too much in

- The `#r4ds` community have [TidyTuesday](#) which makes use of the ideas in Wickham and Grolemund (2017)
- History of visualisation
 - Friendly (2005)
 - Friendly (2006)
- Visualisation to help decision making
 - Tufte (1997)
- ggplot2 resources
 - RStudio (2021)
 - Chang (2017)

References

Chang, Winston. 2017. *R Graphics Cookbook: Practical Recipes for Visualizing Data*. 2nd ed. O'Reilly Media.

<http://www.cookbook-r.com/Graphs/>.

Friendly, M. 2005. "Milestones in the History of Data Visualization: A Case Study in Statistical Historiography." In *Classification: The Ubiquitous Challenge*, edited by C. Weihs and W. Gaul, 34–52. New York: Springer.

<http://www.math.yorku.ca/SCS/Papers/gfkl.pdf>.

"A Brief History of Data Visualization." In *Handbook of Computational Statistics: Data Visualization*, edited by C.

Chen, W. Härdle, and A Unwin. Vol. III. Heidelberg: Springer-Verlag. <http://www.datavis.ca/papers/hbook.pdf>.

Kibirige, Hassan. 2020. *A Grammar of Graphics for Python*. <https://plotnine.readthedocs.io/en/stable/index.html>.

Nolan, Deborah, and Terry P Speed. 2001. *Stat Labs: Mathematical Statistics Through Applications*. Springer Science & Business Media.

Pantoliano, Mike. 2012. "Data Visualization Principles: Lessons from Tufte." 2012. <https://moz.com/blog/data-visualization-principles-lessons-from-tufte>